

SaaS Monitor

SaaS Monitoring Platform

Log Analysis and Monitoring Application
with ELK Stack, NoSQL, and Web Interface

Technical Documentation

Version: 1.0
Date: January 3, 2026
Institution: IT Business School
Scenario: D - SaaS Web Application

Contents

1	Introduction	2
1.1	Context	2
1.2	Problem Statement	2
1.3	Objectives	2
1.4	Chosen Scenario: SaaS Web Application	2
2	Architecture	3
2.1	Global Architecture	3
2.2	Components	3
2.3	Service Ports	3
3	Technical Specifications	4
3.1	Technology Stack	4
3.2	Developed Modules	4
4	Code Excerpts	5
4.1	Prometheus Metrics	5
4.2	Health Check Endpoint	5
5	ELK Configuration	6
5.1	Logstash Pipeline	6
5.2	Elasticsearch Index Mapping	6
6	Data Models	7
6.1	MongoDB Collections	7
6.2	Redis Keys	7
7	Installation Guide	8
7.1	Prerequisites	8
7.2	Quick Start	8
7.3	Verification	8
8	User Guide	9
8.1	Dashboard	9
8.2	File Upload	9
8.3	Search	9
8.4	Live Logs	9

8.5 Grafana Access	9
9 Testing and Validation	10
9.1 API Tests	10
9.2 Validation Checklist	10
10 Challenges and Solutions	11
11 Future Improvements	12
12 Conclusion	13
13 Appendices	14
13.1 API Reference	14
13.2 References	14

Chapter 1

Introduction

1.1 Context

In distributed systems, log production has become exponential. These logs from web applications, microservices, databases, and APIs constitute crucial information for incident detection, performance analysis, security compliance, and business intelligence.

1.2 Problem Statement

Modern enterprises face challenges: logs dispersed across servers, no centralization, manual searching, no automatic alerting, and no visualization of trends and patterns.

1.3 Objectives

This platform must:

1. Centralize log collection from different sources
2. Intelligently index data for ultra-fast search
3. Provide relevant visualizations
4. Offer an intuitive web interface
5. Be easily deployable via containerization

1.4 Chosen Scenario: SaaS Web Application

We selected **Scenario D: SaaS Application** - a monitoring platform for a SaaS used by thousands of client companies.

Types of logs: Web server logs, application logs (Flask), database logs, performance logs, API logs.

KPIs: Error rate by service, average response time, slowest endpoints, errors per hour, active users.

Chapter 2

Architecture

2.1 Global Architecture




Figure 2.1: System Architecture

2.2 Components

Frontend: Bootstrap 5, Chart.js, Socket.IO 7.17, Kibana 8.11

Backend: Flask 3.0, Flask-SocketIO, Gevent

Databases: MongoDB 7, Redis 7

ELK Stack: Elasticsearch 8.11, Logstash

Observability: Prometheus, Grafana, Jaeger

2.3 Service Ports

Service	Port	Service	Port
webapp	5000	kibana	5601
elasticsearch	9200	prometheus	9090
logstash	5044	grafana	3000
mongodb	27017	jaeger	16686
redis	6379		

Table 2.1: Port Mapping

Chapter 3

Technical Specifications

3.1 Technology Stack

Technology	Version	Purpose
Python	3.11	Backend runtime
Flask	3.0.0	Web framework
Flask-SocketIO	5.3.6	Real-time WebSocket
Elasticsearch	8.11.0	Search engine
Logstash	7.17.0	Log processing
MongoDB	7	Document storage
Redis	7	Caching
Prometheus	2.48.0	Metrics collection
Grafana	10.2.2	Dashboards
Docker Compose	2.0+	Orchestration

Table 3.1: Technology Versions

3.2 Developed Modules

File Management: Upload, validation, processing of CSV/JSON logs

Web Interface: Dashboard, search, upload, files, live logs pages

MongoDB Integration: Metadata storage with connection pooling

Real-Time Streaming: WebSocket with Redis message queue

Observability: Prometheus metrics, health checks, structured logging

Chapter 4

Code Excerpts

4.1 Prometheus Metrics

```
1 from prometheus_client import Counter, Histogram, Gauge
2
3 http_requests_total = Counter(
4     'http_requests_total', 'Total HTTP requests',
5     ['method', 'endpoint', 'status_code']
6 )
7
8 http_request_latency_seconds = Histogram(
9     'http_request_latency_seconds', 'Request latency',
10    ['method', 'endpoint'],
11    buckets=[0.01, 0.05, 0.1, 0.25, 0.5, 1.0, 2.5, 5.0]
12 )
```

4.2 Health Check Endpoint

```
1 @app.route('/api/health')
2 def health_check():
3     health = {'status': 'healthy', 'checks': {}, 'system': {}}
4
5     # Elasticsearch check
6     es_start = time.time()
7     if es_client.ping():
8         health['checks']['elasticsearch'] = {
9             'status': 'healthy',
10            'response_time_ms': (time.time() - es_start) * 1000
11        }
12
13     # System metrics
14     health['system'] = {
15         'cpu_percent': psutil.cpu_percent(),
16         'memory_percent': psutil.virtual_memory().percent
17     }
18
19     return jsonify(health)
```

Chapter 5

ELK Configuration

5.1 Logstash Pipeline

```
1 input {
2   tcp { port => 5000; codec => json_lines }
3   file { path => "/data/uploads/*.csv"; tags => ["csv"] }
4 }
5
6 filter {
7   if "csv" in [tags] {
8     csv {
9       columns => ["timestamp", "level", "endpoint",
10                  "status_code", "response_time_ms", "message"]
11     }
12   }
13   date { match => ["timestamp", "ISO8601"] }
14 }
15
16 output {
17   elasticsearch {
18     hosts => ["http://elasticsearch:9200"]
19     index => "saas-logs-%{+YYYY.MM.dd}"
20   }
21 }
```

5.2 Elasticsearch Index Mapping

Field	Type	Purpose
@timestamp	date	Event time
level	keyword	Log level
endpoint	keyword	API path
status_code	integer	HTTP status
response_time_ms	float	Latency
message	text	Log content

Table 5.1: Field Mapping

Chapter 6

Data Models

6.1 MongoDB Collections

```
files: {_id, filename, file_type, file_size, status, uploaded_at, records_count}  
users: {_id, username, email, password_hash, created_at, role}  
search_history: {_id, query, filters, results_count, timestamp, user_id}
```

6.2 Redis Keys

```
cache:stats:*      - Statistics (TTL: 60s)  
cache:search:*     - Search results (TTL: 300s)  
session:user:*    - User sessions  
queue:logs        - Log streaming queue
```

Chapter 7

Installation Guide

7.1 Prerequisites

Docker 20.10+, Docker Compose 2.0+, 8GB RAM minimum.

7.2 Quick Start

```
1 git clone https://github.com/mohamedlandolsi/saas-monitoring-platform.git
2 cd saas-monitoring-platform
3 docker compose up -d
```

7.3 Verification

```
1 docker compose ps          # Check services
2 curl http://localhost:5000/api/health # Test health
3 curl http://localhost:5000/metrics   # Test metrics
```

Chapter 8

User Guide

8.1 Dashboard

Access `http://localhost:5000` to view KPIs: total logs, error rate, average response time, slowest endpoints.

8.2 File Upload

Navigate to Upload page, drag and drop CSV/JSON files, monitor progress, view record count on completion.

8.3 Search

Use search box for keyword search, apply filters (level, time, status code), save frequent searches.

8.4 Live Logs

Access via green “Live” button. Features: pause/resume, level filter, auto-scroll, desktop notifications for errors.

8.5 Grafana Access

URL: `http://localhost:3000`, credentials: admin/admin123. Dashboard: SaaS Monitoring Overview.

Chapter 9

Testing and Validation

9.1 API Tests

Endpoint	P50	P90	P99
/api/health	25ms	45ms	120ms
/api/stats	80ms	150ms	350ms
/api/search	120ms	280ms	650ms

Table 9.1: Response Time Percentiles

9.2 Validation Checklist

- ✓ All API endpoints respond correctly
- ✓ Health checks return accurate status
- ✓ Real-time streaming functional
- ✓ File upload and processing works
- ✓ Search returns relevant results
- ✓ Prometheus metrics exposed
- ✓ Grafana dashboards display data

Chapter 10

Challenges and Solutions

Elasticsearch Memory: Reduced heap to 256MB, added resource limits

Logstash Complexity: Used tags for input identification, conditional filters

WebSocket Scalability: Flask-SocketIO with Redis message queue

Kibana Dependencies: Docker health checks with service _healthy condition

Performance: Redis caching, query optimization, pagination

Chapter 11

Future Improvements

Short-term:

- ML anomaly detection
- Slack/email alerting
- Custom dashboards

Long-term:

- Kubernetes deployment
- Multi-tenancy support
- API versioning

Chapter 12

Conclusion

This project successfully delivered a comprehensive SaaS Log Monitoring Platform integrating ELK Stack, MongoDB, Redis, and modern observability tools.

Achievements:

- Centralized log collection from multiple sources
- Sub-second search with Elasticsearch
- Real-time streaming via WebSocket
- Prometheus metrics and Grafana dashboards
- Containerized deployment with Docker Compose

The modular architecture ensures extensibility and easy scaling.

Chapter 13

Appendices

13.1 API Reference

Endpoint	Method	Description
/api/health	GET	Service health
/api/stats	GET	Log statistics
/metrics	GET	Prometheus metrics
/api/logs	GET	Query logs
/api/search	GET	Search logs
/api/upload	POST	Upload file
/api/files	GET	List files

Table 13.1: API Endpoints

13.2 References

1. Elasticsearch: <https://www.elastic.co/guide/>
2. Flask: <https://flask.palletsprojects.com/>
3. Docker: <https://docs.docker.com/>
4. Prometheus: <https://prometheus.io/docs/>